











LLaVA-NeXT-7B



LLaVA-Video-7B



Baseline FastV CDPPruner VScan STAR-Pro (Ours)

Efficiency Studies — latency, compute, peak GPU memory, speedup, and relative performance

LLaVA-NeXT-7B uses retained-token budgets; LLaVA-Video-7B uses tokens per frame over 64 frames.